

UPPSALA UNIVERSITY

DEEP LEARNING

---

# Hand-in assignment (3)

---

Tong You

June 10, 2021



# Contents

<b>1</b>	<b>Introduction</b>	<b>2</b>
<b>2</b>	<b>The dataset</b>	<b>2</b>
<b>3</b>	<b>A simple language model</b>	<b>3</b>
3.1	Results . . . . .	4
<b>4</b>	<b>Improving the recurrent model</b>	<b>5</b>
4.1	Results - gradient clipping . . . . .	5
4.2	Results - LSTM . . . . .	6
4.3	Results - additional modifications . . . . .	8
<b>5</b>	<b>Appendix</b>	<b>10</b>

# 1 Introduction

Whereas in the previous two assignments we were working with image data (MNIST and Warwick data), in this final assignment we will implement a language model. We will be using the [PennTreebank](#) (PTB) as our dataset. This is a collection of articles from the Wall Street Journal. We will implement the language model in several stages. First, we need to preprocess the data. Next we will implement a (simple) Elmann RNN. Finally, we will look at several different ways to improve the model. We will show generated example phrases at each stage in order to show the differences. The model will be implemented in Pytorch using an RTX 2080 Ti in order to do GPU training.

## 2 The dataset

As mentioned previously, we will be using the PTB dataset. This is split up into three splits: a training split, validation split, and testing split. Before we can feed this data into our language model we need to preprocess the data, generate a dictionary with our keys being the unique words in the three splits and the values being an unique integer corresponding to a particular unique word. Before we proceed, we will be showing some example sentences. We pick two random sentences from each split.

- Example training set phrase: in addition the cray-3 will contain N processors twice as many as the largest current supercomputer
- Example training set phrase: what 's more the test and learning materials are both produced by the same company <unk> a joint venture of mcgraw-hill inc. and macmillan 's parent britain 's maxwell communication corp
- Example validation set phrase: the ghost of the soviet <unk> discovered in cuba back in the <unk> costs just a few hundred million the price of the caribbean command in key west that president carter created in N
- Example validation set phrase: mr. wolf has <unk> merger advice from a major wall street securities firm relying instead only on a takeover lawyer peter <unk> of <unk> <unk> slate <unk> & flom

- Example test set phrase: the executive denied speculation that saatchi was bringing in the new chief executive officer only to clean up the company financially so that the brothers could lead a buy-back
- Example test set phrase: quantum 's lot is mostly tied to polyethylene <unk> used to make garbage bags milk <unk> <unk> toys and meat packaging among other items

The next step is to add an additional term <eos> that denotes the end of a sentence. All code will be attached in the Appendix. Next we want to check the size of the three splits. For this count we consider the special characters <eos> and <unk>. The size of the splits:

- Number of training words: 929589
- Number of validation words: 73760
- Number of testing words: 82430

From the above we want to generate a dictionary as mentioned earlier. In order to generate the dictionary we first flatten the three splits and combine the splits into one big list. Then we use the set function in Python. This will generate a set with the unique words. From this we can easily see how many unique words there are. There are 10000 unique words in the entire PTB dataset. Finally, we use this dictionary to map our words to a unique integer. To generate sample phrases we would look up the word corresponding to any unique integer. Our input to the RNN will therefore consist of a vector of integers. Next we will look at the language model.

### 3 A simple language model

In the previous section we did some preprocessing on the PTB dataset. Now we are ready to implement a simple RNN using PyTorch. For implementation details we refer to the Appendix, however code that is deemed important is explained in the following text. After some initial tweaking, an "ideal" set of hyperparameters was chosen. The test perplexity still doesn't reach the value of a highly-optimized Elman RNN. However, we will implement some methods to try to lower the perplexity. In order to make an one-to-one comparison between the sections we will use the same input sentence(s) when generating sample phrases.

### 3.1 Results

To start with, we will show results of a vanilla RNN without any gradient clipping. In order to highlight differences between different methods we will try to use the same parameters in order to really show the gains that were made. With normal SGD it was very difficult get a low perplexity that starts to approach a well-tuned Elmann RNN. The following loss curves correspond to a training batch size of 32 and validation/testing batch sizes of 512. The sequence length for all splits was 50. Both the hidden and embedding dimensions were set to 500. The RNN was implemented with a layer number of 1 first. The loss was averaged each epoch to yield less noisier training/validation curves. We will also use dropout since that really improved the training for the simple RNN (value of 0.5). To start with, we will show results of a vanilla RNN without any gradient clipping. In order to highlight differences between different methods we will try to use the same parameters in order to really show the gains that were made. With normal SGD it was very difficult get a low perplexity that starts to approach a well-tuned Elmann RNN. The following loss curves correspond to a training batch size of 32 and validation/testing batch sizes of 512. The sequence length for all splits was 50. Both the hidden and embedding dimensions were set to 500. The RNN was implemented with a layer number of 1 first. The loss was averaged each epoch to yield less noisier training/validation curves. We will also use dropout since that really improved the training for the simple RNN (value of 0.5). The learning rate was 0.5. The loss curve below shows the results for the simple RNN. This model was then evaluated on the test split and we got a perplexity of 214.15121520493514. This is still quite high, but not as high as initial testing of parameters.




Figure 1: Training and validation loss averaged over all iterations each epoch. The final losses are annotated in the figure. The RNN has one layer.

We will see in the next section how we can improve this.

## 4 Improving the recurrent model

As mentioned earlier, there are several ways to improve upon the previous result. In addition we will also show examples of generated phrases by providing a seeding input to the trained model.

### 4.1 Results - gradient clipping

We will first show some results of gradient clipping applied to a two-layer RNN. We will use the same parameters, but after some preliminary testing we found that we can use quite high learning rates, so we chose 10. We clipped the gradients at 0.5 using the L-2 norm. The one-layer RNN got higher perplexity than the one-layer network with perplexity in testing so we decided to show results for a two-layer network with gradient clipping.

This network gets higher test perplexity than the above network. The test perplexity after 50 epochs was 247.3608380295578.




Figure 2: Training and validation loss averaged over all iterations each epoch. Training was done with gradient clipping. The final losses are annotated in the figure. The RNN has two layers.

Next we will show a generated phrase giving an input sequence. We will use a greedy method. However, the current model outputs the same word when we specify how many words we want to generate. Several attempts to fix this didn't work, so only one generated phrase is shown. However, it seems that the model outputs the most frequent word. The seeding sequence is 2 words long and is created by using uniformly generated random integers up to the number of unique words. We started with "there" and "unrelated".

- Generated phrase: "there" "unrelated" "with" "with" "with"

## 4.2 Results - LSTM

Since the Elmann RNN is quite "simple", we will next use an LSTM. These are much less prone to exploding gradients, although we didn't encounter

such problems for our implementation. We will use the exact same parameters as the results in the subsection above, only changing the RNN to an LSTM. For ease of reading, we will give them again. For the training we used a batch size of 32 and for validation/testing we evaluated on data with batch sizes of 512. The sequences were of length 50 for all splits. We also used the same embedding and hidden dimensions of 500. And the LSTM in this case has two layers. We also found that a dropout of 0.8 yielded quite good results. The learning rate was set to 10.0 with gradient clipping with a max L2-norm of 0.5. We also trained for 50 epochs using standard SGD. The final test perplexity was 159.23738342645728.




Figure 3: Training and validation loss averaged over all iterations each epoch. Training was done with gradient clipping. The final losses are annotated in the figure. Instead of a two-layer RNN, a two-layer LSTM was trained.

Next we show an example of a generated phrase. We used as our input sequence: "there" and "competitiveness". We wanted to generate three words.

- Generated phrase: "there" "competitiveness" "is" "<unk>" "<unk>" .



### 4.3 Results - additional modifications

There are different changes we can apply to our current model and see if we can spot any improvements. We will use the same settings as used above, but swap the LSTM with an GRU. Training will be done in the same way.




Figure 4: Training and validation loss averaged over all iterations each epoch. Training was done with gradient clipping. The final losses are annotated in the figure. A two-layer GRU was used to train the PTB dataset.

The final test perplexity was slightly lower than the LSTM: 157.50157610436733. An example of a generated phrase:

- Generated phrase: "there" "unknown" "is" "a" "<unk>" .

The trained model was seeded with a sequence in a similar way as above. The model generates much more promising results, however it still outputs the same word after a few generations. We will finally try to improve on the previous model by using a custom learning rate scheduler. We will still use the GRU but instead use a cosine learning rate scheduler and train for 50 epochs again. We will use SGD again with a learning rate of 10, but this

will be modified by the cosine learning rate scheduler. The scheduler was implemented as follows:

```
1 cosine_sched = optim.lr_scheduler.CosineAnnealingLR(optimizer,
  ↳ T_max = 500, eta_min = 0.1, last_epoch = -1, verbose =
  ↳ True)
```




Figure 5: Training and validation loss averaged over all iterations each epoch. Training was done with gradient clipping. The final losses are annotated in the figure. A two-layer GRU was used to train the PTB dataset. This time a cosine learning rate scheduler was used.

With the same parameters as the GRU with SGD and no learning rate scheduling, we now get a final perplexity of 146.02419039272354. This is about an 11 point difference compared to the GRU without rate scheduling. We can see in the loss curves that the loss peaked a bit after about 25 epochs, but decreased again afterwards. Better results might be expected by combining an exponential decay scheduler with a cosine learning rate scheduler. Despite the improvements in perplexity using the same seeding mechanism

as we had previously the generated phrases still repeat the same word after only a few words:

- Generated phrase: "there" "brewing" "is" "a" "<unk>".

## 5 Appendix

Since the code was implemented in a single Jupyter Notebook, the code in the different cells will be attached here as a single file.

```
1 import torch
2 import torch.nn as nn
3 import torch.optim as optim
4 import torch.nn.functional as F
5
6 import functools
7 import operator
8 import random
9 import pickle
10
11 import numpy as np
12
13 from sklearn.metrics import confusion_matrix
14 from torch.utils.data import TensorDataset
15 from torch.utils.data import DataLoader
16 from matplotlib import pyplot as plt
17
18 # Exercise 1
19 # Training data
20 keepNew = True
21 with open('PTB/ptb.train.txt', mode = 'r', newline = None) as
    ↪ train_f:
22     train_dat = train_f.read().splitlines(keepNew)
23
24 # Validation data
25 with open('PTB/ptb.valid.txt', mode = 'r', newline = None) as
    ↪ valid_f:
26     valid_dat = valid_f.read().splitlines(keepNew)
```

```

27
28 # Testing data
29 with open('PTB/ptb.test.txt', mode = 'r', newline = None) as
    ↪ test_f:
30     test_dat = test_f.read().splitlines(keepNew)
31
32 # Printing examples phrases
33 n_examp = np.random.choice(1000, size = 6, replace = False)
34
35 print("Example training set phrase:", train_dat[n_examp[0]])
36 print("Example training set phrase:", train_dat[n_examp[1]])
37 print("Example validation set phrase:", valid_dat[n_examp[2]])
38 print("Example validation set phrase:", valid_dat[n_examp[3]])
39 print("Example test set phrase:", test_dat[n_examp[4]])
40 print("Example test set phrase:", test_dat[n_examp[5]])
41
42 # Adding <eos>
43 train_proc, valid_proc, test_proc = [], [], []
44 for l_train in train_dat:
45     l_mod = l_train.replace('\n', '<eos>')
46     train_proc.append(l_mod)
47
48 for l_valid in valid_dat:
49     l_mod = l_valid.replace('\n', '<eos>')
50     valid_proc.append(l_mod)
51
52 for l_train in test_dat:
53     l_mod = l_train.replace('\n', '<eos>')
54     test_proc.append(l_mod)
55
56 # Split up each line in individual words
57 train_words, valid_words, test_words = [], [], []
58 for tp in train_proc:
59     train_words.append(tp.split())
60
61 for tp in valid_proc:
62     valid_words.append(tp.split())
63

```

```

64 for tp in test_proc:
65     test_words.append(tp.split())
66
67 # Flatten list of lists into a single list
68 train_words = functools.reduce(operator.iconcat, train_words,
    ↪ [])
69 valid_words = functools.reduce(operator.iconcat, valid_words,
    ↪ [])
70 test_words = functools.reduce(operator.iconcat, test_words,
    ↪ [])
71
72 num_train = len(train_words)
73 num_valid = len(valid_words)
74 num_test = len(test_words)
75
76 print("Number of training words:", num_train)
77 print("Number of validation words:", num_valid)
78 print("Number of testing words:", num_test)
79
80 # Building a dictionary
81 all_words = train_words + valid_words + test_words
82 set_words = set(all_words)
83 num_unique = len(set_words)
84
85 print("Number of unique words in training + validation +
    ↪ testing splits:", num_unique)
86
87 num_id = np.random.choice(num_unique, size = num_unique,
    ↪ replace = False)
88
89 n = 0
90 unique_dict = {}
91 for uw in set_words:
92     unique_dict.update({uw : num_id[n]})
93     n += 1
94
95 # Replacing all words in the training/validation/testing
    ↪ splits with their integer representation

```

```

96 train_ints, valid_ints, test_ints = [], [], []
97
98 for word in train_words:
99     int_rep = unique_dict[word]
100     train_ints.append(int_rep)
101
102 for word in valid_words:
103     int_rep = unique_dict[word]
104     valid_ints.append(int_rep)
105
106 for word in test_words:
107     int_rep = unique_dict[word]
108     test_ints.append(int_rep)
109
110 # Check if CUDA is available
111 device = torch.device("cuda") if torch.cuda.is_available()
112     ↪ else torch.device("cpu")
113 print("Device:", torch.cuda.get_device_name(device))
114
115 # Resetting model function
116 # Credits:
117     ↪ https://discuss.pytorch.org/t/reset-model-weights/19180/4
118 def reset_model(model):
119     for layer in model.children():
120         if hasattr(layer, 'reset_parameters'):
121             layer.reset_parameters()
122
123 # Exercise 2
124 # Convert our training/validation/testing splits to Torch
125     ↪ tensors
126 train_dat = torch.tensor(train_ints)
127 valid_dat = torch.tensor(valid_ints)
128 test_dat = torch.tensor(test_ints)
129
130 ### Data loading
131 batch_size = 32
132 batch_eval = 512
133 batch_test = 512

```

```

131 seq_train = 50
132 seq_valid = 50
133 seq_test = 50
134
135 s_train_l = num_train // seq_train
136 s_valid_l = num_valid // seq_valid
137 s_test_l = num_test // seq_test
138
139 # Trim training/validation/testing data and reshape into
140 ↪ tensor of num_sequences by sequence_length
141 # Training data
142 train_seq = torch.narrow(train_dat, 0, 0, seq_train *
143 ↪ s_train_l)
144 train_lab = torch.roll(train_seq, shifts = -1, dims = 0)
145
146 train_seq = train_seq.reshape(s_train_l, seq_train)
147 train_lab_seq = train_lab.reshape(s_train_l, seq_train)
148
149 # Validation data
150 valid_seq = torch.narrow(valid_dat, 0, 0, seq_valid *
151 ↪ s_valid_l)
152 valid_lab = torch.roll(valid_seq, shifts = -1, dims = 0)
153
154 valid_seq = valid_seq.reshape(s_valid_l, seq_valid)
155 valid_lab_seq = valid_lab.reshape(s_valid_l, seq_valid)
156
157 # Testing data
158 test_seq = torch.narrow(test_dat, 0, 0, seq_test * s_test_l)
159 test_lab = torch.roll(test_seq, shifts = -1, dims = 0)
160
161 test_seq = test_seq.reshape(s_test_l, seq_test)
162 test_lab_seq = test_lab.reshape(s_test_l, seq_test)
163
164 # Divide training and validation data into correct
165 ↪ mini-batches
166 num_batches = train_seq.shape[0] // batch_size
167 valid_batches = valid_seq.shape[0] // batch_eval
168 test_batches = test_seq.shape[0] // batch_test

```

```

165
166 training_set = TensorDataset(train_seq.to(device),
    ↪ train_lab_seq.to(device))
167 training_loader = DataLoader(training_set, shuffle = False,
    ↪ batch_size = num_batches)
168
169 valid_set = TensorDataset(valid_seq.to(device),
    ↪ valid_lab_seq.to(device))
170 valid_loader = DataLoader(valid_set, shuffle = False,
    ↪ batch_size = valid_batches)
171
172 test_set = TensorDataset(test_seq.to(device),
    ↪ test_lab_seq.to(device))
173 test_loader = DataLoader(test_set, shuffle = False, batch_size
    ↪ = test_batches)
174
175 ### RNN code
176 # Embedding parameter
177 embed_dim = 500
178
179 # RNN parameters
180 hidden_dim = 500
181 in_size = embed_dim
182 n_layers = 2
183
184 # Vanilla RNN
185 class ElmanRNN(nn.Module):
186     def __init__(self, input_size, hidden_size, num_layers,
    ↪ num_embeddings, embedding_dim, hidden_dim, num_unique,
    ↪ drop_out):
187         super(ElmanRNN, self).__init__()
188
189         self.input_size = input_size
190         self.hidden_size = hidden_size
191         self.num_layer = num_layers
192         self.drop_out = drop_out
193         self.num_embeddings = num_embeddings
194         self.embedding_dim = embedding_dim

```



```

195
196     self.embed = nn.Embedding(num_embeddings,
    ↪     embedding_dim)
197     self.elman = nn.RNN(input_size, hidden_size,
    ↪     num_layers, dropout = drop_out, batch_first =
    ↪     True)
198     self.linear = nn.Linear(hidden_dim, num_unique)
199
200     def forward(self, mod_input):
201         word_embed = self.embed(mod_input)
202         rnn_out, hidden_out = self.elman(word_embed)
203         rnn_out = self.linear(rnn_out)
204         rnn_out = rnn_out.view(-1, num_unique)
205
206         return rnn_out, hidden_out
207
208     # LSTM
209     class LSTMRNN(nn.Module):
210         def __init__(self, input_size, hidden_size, num_layers,
    ↪         num_embeddings, embedding_dim, hidden_dim, num_unique,
    ↪         drop_out):
211             super(LSTMRNN, self).__init__()
212
213             self.input_size = input_size
214             self.hidden_size = hidden_size
215             self.num_layer = num_layers
216             self.drop_out = drop_out
217             self.num_embeddings = num_embeddings
218             self.embedding_dim = embedding_dim
219
220             self.embed = nn.Embedding(num_embeddings,
    ↪             embedding_dim)
221             self.lstm = nn.LSTM(input_size, hidden_size,
    ↪             num_layers, dropout = drop_out, batch_first =
    ↪             True, bidirectional = False)
222             self.linear = nn.Linear(hidden_dim, num_unique)
223
224     def forward(self, mod_input):

```

```

225         word_embed = self.embed(mod_input)
226         lstm_out, hidden_out = self.lstm(word_embed)
227         lstm_out = self.linear(lstm_out)
228         lstm_out = lstm_out.view(-1, num_unique)
229
230         return lstm_out, hidden_out
231
232     # GRU
233     class GRURNN(nn.Module):
234         def __init__(self, input_size, hidden_size, num_layers,
235             ↪ num_embeddings, embedding_dim, hidden_dim, num_unique,
236             ↪ drop_out):
237             super(GRURNN, self).__init__()
238
239             self.input_size = input_size
240             self.hidden_size = hidden_size
241             self.num_layer = num_layers
242             self.drop_out = drop_out
243             self.num_embeddings = num_embeddings
244             self.embedding_dim = embedding_dim
245
246             self.embed = nn.Embedding(num_embeddings,
247             ↪ embedding_dim)
248             self.gru = nn.GRU(input_size, hidden_size, num_layers,
249             ↪ dropout = drop_out, batch_first = True,
250             ↪ bidirectional = False)
251             self.linear = nn.Linear(hidden_dim, num_unique)
252
253         def forward(self, mod_input):
254             word_embed = self.embed(mod_input)
255             gru_out, hidden_out = self.gru(word_embed)
256             gru_out = self.linear(gru_out)
257             gru_out = gru_out.view(-1, num_unique)
258
259             return gru_out, hidden_out
260
261     elman_rnn = ElmanRNN(in_size, hidden_dim, n_layers,
262         ↪ num_unique, embed_dim, hidden_dim, num_unique, 0.5)

```

```

257 lstm_rnn = LSTMRNN(in_size, hidden_dim, n_layers, num_unique,
    ↪ embed_dim, hidden_dim, num_unique, 0.8)
258 gru_rnn = GRURNN(in_size, hidden_dim, n_layers, num_unique,
    ↪ embed_dim, hidden_dim, num_unique, 0.8)
259
260 # Selecting model
261 model = gru_rnn
262
263 reset_model(model)
264 model.to(device)
265
266 # Optimizer
267 l_rate = 10.0
268 sgd = optim.SGD(model.parameters(), lr = l_rate, weight_decay
    ↪ = 0, momentum = 0.0)
269 adam = optim.Adam(model.parameters(), lr = l_rate, betas =
    ↪ (0.9, 0.999), eps = 1e-08, weight_decay = 0.0, amsgrad =
    ↪ False)
270
271 optimizer = sgd
272
273 # Cross entropy loss
274 loss = nn.CrossEntropyLoss()
275
276 # Training parameters
277 num_epochs = 50
278 train_epoch = np.zeros(num_epochs)
279 valid_epoch = np.zeros(num_epochs)
280
281 train_loss = []
282 valid_loss = []
283 test_loss = []
284
285 # Gradient clipping
286 clipGrad = True
287
288 # Learning rate scheduler

```

```

289 exp_sched = optim.lr_scheduler.ExponentialLR(optimizer, gamma
    ↪ = 0.5, last_epoch = -1, verbose = True)
290 cosine_sched = optim.lr_scheduler.CosineAnnealingLR(optimizer,
    ↪ T_max = 500, eta_min = 0.1, last_epoch = -1, verbose =
    ↪ True)
291 #plat_sched =
    ↪ optim.lr_scheduler.ReduceLROnPlateau(optimizer, mode =
    ↪ 'min', factor = 0.1, patience = 10, threshold = 0.0001,
    ↪ threshold_mode = 'rel', cooldown = 0, min_lr = 0, eps =
    ↪ 1e-08, verbose = True)
292
293 sched = cosine_sched
294 useScheduler = True
295
296 # Training and evaluation on validation data
297 for epoch in range(num_epochs):
298     train_lossss = []
299     valid_lossss = []
300
301     for data, lab_train in training_loader:
302         model.train()
303         optimizer.zero_grad()
304
305         pred_out, hid_out = model(data)
306
307         ce_loss = loss(pred_out, lab_train.flatten())
308
309         train_loss.append(ce_loss.item())
310         train_lossss.append(ce_loss.item())
311
312         ce_loss.backward()
313
314         # Using gradient clipping
315         if clipGrad:
316             nn.utils.clip_grad_norm_(model.parameters(),
    ↪ max_norm = 0.5, norm_type = 2.0)
317
318     optimizer.step()

```

```

319
320     # Using learning rate scheduler
321     if useScheduler:
322         sched.step()
323
324     train_epoch[epoch] = np.mean(train_lossss)
325
326     for data_eval, lab_eval in valid_loader:
327         model.eval()
328         with torch.no_grad():
329             valid_preds, valid_hid = model(data_eval)
330
331             ce_valid = loss(valid_preds, lab_eval.flatten())
332             valid_loss.append(ce_valid.item())
333             valid_lossss.append(ce_valid.item())
334
335     valid_epoch[epoch] = np.mean(valid_lossss)
336
337     print("Epoch: %s" % (epoch + 1))
338
339     # Evaluation on test data
340     model.eval()
341     for data_test, lab_test in test_loader:
342         with torch.no_grad():
343             test_preds, test_hid = model(data_test)
344
345             ce_test = loss(test_preds, lab_test.flatten())
346             test_loss.append(ce_test.item())
347
348     plt.figure(1, figsize = (6.4, 4.8))
349     train, = plt.plot(train_epoch, 'r')
350     valid, = plt.plot(valid_epoch, 'b')
351     plt.xlabel("Epoch")
352     plt.ylabel("Loss")
353     plt.title("Averaged training and validation loss")
354     plt.legend([train, valid], ['Train loss', 'Validation loss'])

```

```

355 plt.annotate("Final average training loss: %s" %
    ↪ (train_epoch[-1]) ,xycoords = 'figure fraction', xy =
    ↪ (0.25,0.55))
356 plt.annotate("Final average validation loss: %s" %
    ↪ (valid_epoch[-1]), xycoords = 'figure fraction', xy =
    ↪ (0.25,0.50))
357 plt.savefig("train_valid_epoch", dpi = 500)
358 print("Final average training loss: %s." % train_epoch[-1])
359 print("Final average validation loss: %s." % valid_epoch[-1])
360 print()
361
362 plt.figure(2, figsize = (6.4, 4.8))
363 train_it, = plt.plot(train_loss, 'r')
364 valid_it, = plt.plot(valid_loss, 'b')
365 plt.xlabel("Iteration")
366 plt.ylabel("Loss")
367 plt.title("Training and validation loss")
368 plt.legend([train_it, valid_it], ['Train loss', 'Validation
    ↪ loss'])
369 plt.annotate("Final training loss: %s" % (train_loss[-1])
    ↪ ,xycoords = 'figure fraction', xy = (0.25,0.55))
370 plt.annotate("Final validation loss: %s" % (valid_loss[-1]),
    ↪ xycoords = 'figure fraction', xy = (0.25,0.50))
371 #plt.savefig("train_valid_loss_iter", dpi = 500)
372 print("Final training loss: %s." % train_loss[-1])
373 print("Final validation loss: %s." % valid_loss[-1])
374 print()
375
376 plt.figure(3, figsize = (6.4, 4.8))
377 train_fin_ep, = plt.plot(train_lossss, 'r')
378 valid_fin_ep, = plt.plot(valid_lossss, 'b')
379 plt.xlabel("Iteration")
380 plt.ylabel("Loss")
381 plt.title("Training and validation loss final epoch")
382 plt.legend([train_fin_ep, valid_fin_ep], ['Train loss',
    ↪ 'Validation loss'])

```

```

383 plt.annotate("Training loss final epoch: %s" %
    ↪ (train_lossss[-1]) ,xycoords = 'figure fraction', xy =
    ↪ (0.25,0.55))
384 plt.annotate("Validation loss final epoch: %s" %
    ↪ (valid_lossss[-1]), xycoords = 'figure fraction', xy =
    ↪ (0.25,0.50))
385 #plt.savefig("train_valid_final_epoch", dpi = 500)
386
387 perplexity = np.exp(np.mean(test_loss))
388 print("Test perplexity: %s." % (perplexity))
389
390 # Generate sample phrases
391 dictionary = unique_dict # number of unique words/tokens is
    ↪ 10000
392 dict_keys = list(dictionary.keys())
393 dict_values = list(dictionary.values())
394
395 # Words
396 w_1 = "there"
397 i_1 = dictionary[w_1]
398 w_2 = "is"
399 i_2 = dictionary[w_2]
400 w_3 = "a"
401 i_3 = dictionary[w_3]
402 w_4 = "why"
403 i_4 = dictionary[w_4]
404 w_5 = "on"
405 i_5 = dictionary[w_5]
406
407 # Input sequence
408 input_tens = torch.randint(num_unique, (2,
    ↪ 1)).long().to(device)
409 input_tens[0] = i_1
410 print("Input", input_tens)
411
412 inputs_list = input_tens.tolist()
413 print("List input", inputs_list)
414

```

```

415 softmax_layer = nn.Softmax(dim = 1)
416 num_words = 5 # how many words to predict
417 model.eval()
418 with torch.no_grad():
419     for n in range(num_words):
420         out, hidden = model(input_tens)
421         probs = softmax_layer(out)
422
423         maxv, maxi = torch.max(probs, dim = 1)
424         sortedv, sortedi = -np.sort(-maxv.cpu().numpy()),
         ↪ -np.sort(-maxi.cpu().numpy())
425
426         next_word = sortedi[0]
427
428         inputs_list.append([next_word])
429
430         input_tens = input_tens.tolist()
431         input_tens.append([next_word])
432         input_tens =
         ↪ torch.tensor(input_tens).long().to(device)
433
434
435 print(inputs_list)
436 output_sentence = []
437
438 for i in inputs_list:
439     word_i = dict_keys[dict_values.index(i[0])]
440     output_sentence.append(word_i)
441
442 print(output_sentence)

```